



Deep Learning with CUDA

Dedicated architectures – CNN


Tomasz Szumlak

WFiIS AGH
27/03/2024, Kraków

CNN – computer vision master




- ❑ One of the most sophisticated ANN models to date is the CNNs or Convolutional Neural Networks
- ❑ The main difference between the general dense model, which we have considered so far, **and CNN is specialisation** – enter the Conv layers
- ❑ Conv layers are strongly motivated by the biological retina





3

Conv layer demo...




<https://cs231n.github.io/convolutional-networks/>

Conv layer properties



- ❑ A few things to note...
 - ❑ We can process **colourful data** (R, G, B channels)
 - ❑ There is **extra padding** introduced around the original picture to get the equal attention of the scanning kernel to each pixel
 - ❑ **Kernels are a bit of an analogue of „advanced” neurons** (kernel parameters are trainable, and they are considered to be model weights)
 - ❑ For the example above, one can consider the kernel to be a 3D object with a single bias weight



Conv layer properties



- ❑ A few things to note...
 - ❑ The size of the kernel and the way it moves across the network architecture can define the input data
 - ❑ The pace of the movement is called the stride (could be 1, 2...)
 - ❑ In our example, we have R-, B- and G-filters representing **one super neuron**
 - ❑ At each scan point, we **calculate the Hadamard product** (that is just an analogue of the grand neuron equation!!)
 - ❑ The total output of the first kernel filter is just a sum of all products plus the bias
 - ❑ In this way, **we create a compressed representation** of the input data
 - ❑ And we can have multiple kernels in each layer



Conv layer properties


- ☐ Kernels represent some spatial features we are looking for
- ☐ On deeper layers, the features get more and more abstract
- ☐ The point is that in the output, the **significant positive** value will indicate the **presence of the feature**, and the **significant negative** will say that there is a **lack of one**
- ☐ **Important statement – Conv layers are designed to detect pattern/spatial features in position invariant way**
- ☐ The layers are faithful to the spatial structure of an input image
- ☐ This makes them robust against moving the picture, resizing, etc.

Conv layer properties




- ❑ Each layer can have many filters/kernels
- ❑ Each of which makes a representation of the input data
- ❑ If we compare the regular dense networks with CNNs, we see that the number of parameters required by CNNs are much, much lower
- ❑ On top of that, CNNs are going to be superior to any dense architecture in computer vision applications

Conv layer properties




Each feature – one kernel

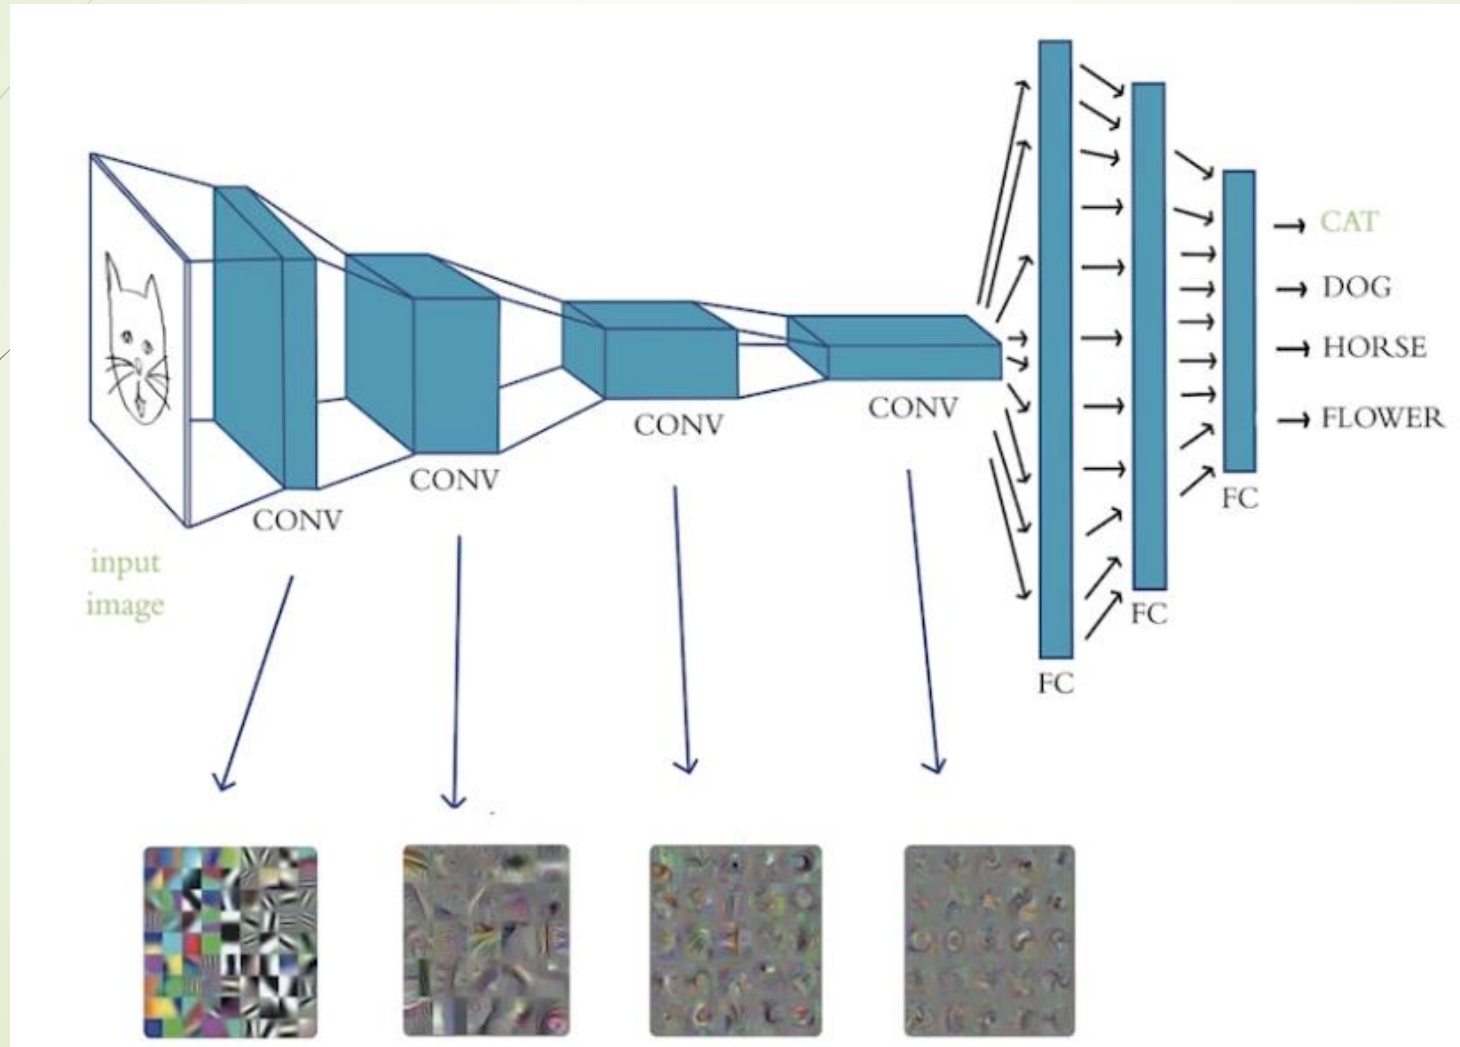


CNN


Many layers – building complex image



CNN - AlexNet




Pooling layer – good partner for CL



It reduces the size of output activations but preserves the depths (MaxPooling)


Our playground – LeNet-like model



Deep models and thir applications




CLASSIFICATION




“BALLOONS”


OBJECT
DETECTION



SEMANTIC
SEGMENTATION



INSTANCE
SEGMENTATION



Capsule network – positional info



Hyper-parameter tuning



- ☐ Perform your model initialisation wisely...
- ☐ Define your cost function according to the task
- ☐ Set the limit – the goal/target, the random chance is the boundary
- ☐ Layers (types, number, sizes)
- ☐ Use monitoring tools to check for the over-fitting
- ☐ Learning rate tuning
- ☐ Batch size tuning

- ☐ Automatic tools (Spearmin, Hyperas, Hyperopt)

YOLO – You Only Look Once



❑ Let's appreciate this for a moment

<https://arxiv.org/abs/1506.02640>



OpenCV – Computer Vision




- ❑ For CV we are mostly interested in manipulating of images, OpenCV is one of the most common and loved free piece of software
- ❑ It is cross-platform – fairly easy to switch between different Oses
- ❑ Usually need to be installed separately
- ❑ If you have anaconda on your system you can go for „**pip install opencv-python**”





```
#import the libraries
import cv2

#printing the version
print(cv2.__version__)
```

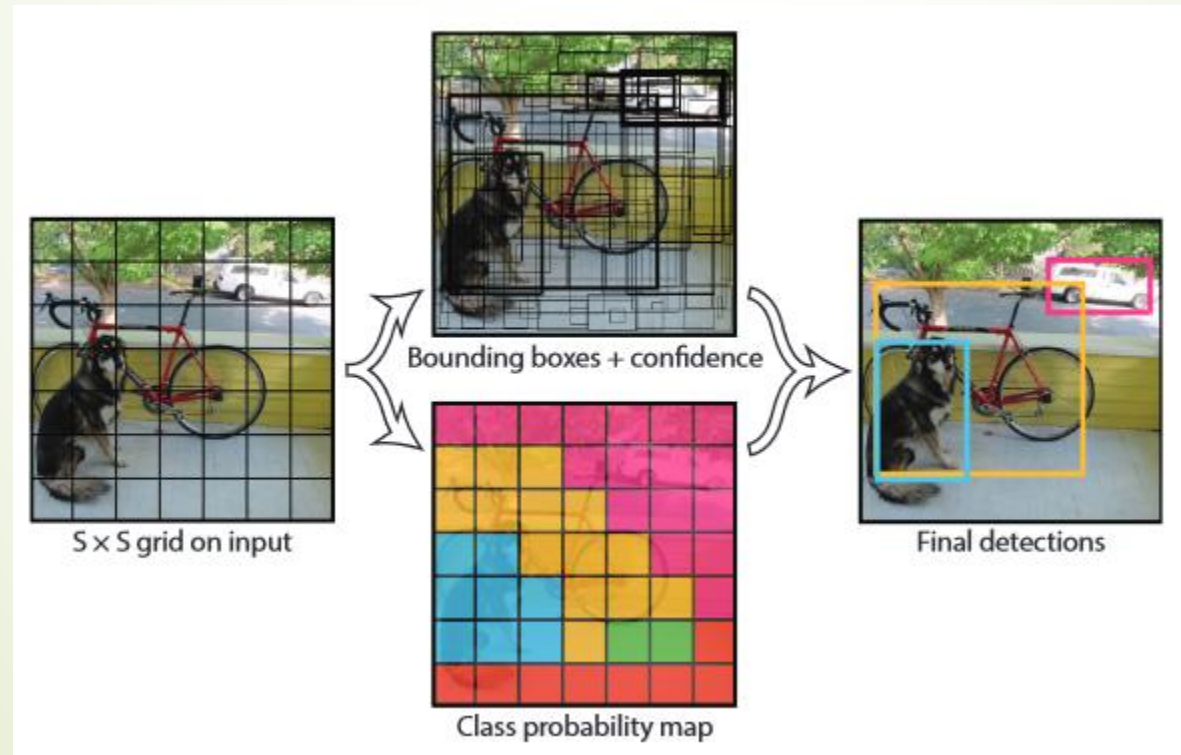

Yolo architecture (2016)



Checking on art datasets



Principle of YOLO




Darknet




- ❑ Original, native environment based on C/CUDA implementation
- ❑ We can make it work on colab-like distributed platforms, this is what we are going to do



YOLO innerworkings





„Virtual patching”




Output feature map from CNN

YOLO innerworkings



Formula for a single bounding box for a given patch

Intersection over union



The end